

Lagrangian Approaches for Knowledge Injection

When soft constraints are enough

Constrained Machine Learning

Let's consider ML problem with **constrained output**

In particular, let's focus on problems in the form:

$$\operatorname{argmin}_{\theta} \{ L(y) \mid y = f(\hat{x}, \omega), g(y) \leq 0 \}$$

Where:

- L is the loss
- \hat{x} is the training input
- y is the ML model output, i.e. $f(x, \omega)$
- ω is the parameter vector (we assume a parameterized model)
- g is a constraint function

Equality csts. can be viewed as double-inequalities (but they admit simplifications)

Constrained Machine Learning

Example 1: logical rules

E.g. hierarchies in multi-class classification ("A dog is also an animal"):

$$y_{i,dog} \leq y_{i,animal}$$

- This constraint is defined over individual examples

Example 2: shape constraints

E.g. input x_j cannot cause the output to decrease (monotonicity)

$$y_i \leq y_k \quad \forall i, k : x_{i,j} \leq x_{k,j} \wedge x_{i,h} = x_{k,h} \forall h \neq j$$

- This is a relational constraint, i.e. defined over multiple examples

Lagrangian Relaxations for Constrained ML

One way to deal with this problem is to rely on a **Lagrangian Relaxation**

Main idea: we turn the constraints into **penalty terms**:

$$\operatorname{argmin}_{\theta} \{ L(y) + \lambda^T \max(0, g(y)) \mid y = f(\hat{x}, \omega) \}$$

Or, alternatively:

$$\operatorname{argmin}_{\theta} \{ L(y) + \lambda^T \max(0, g(y))^2 \mid y = f(\hat{x}, \omega) \}$$

- We use a vector of **multipliers** λ to weight the constraint violations

This is a popular approach for ML with constraints

- One of the first occurrences as Semantic Based Regularization
- The constraints are "distilled" in the model parameters

Lagrangian Relaxations for Constrained ML

Equality constraints can be treated as double inequalities

...Or via the simplified form:

$$\arg \min_{\theta} \{ L(\mathbf{y}) + \lambda^T g(\mathbf{y})^2 \} \quad \text{with: } \mathbf{y} = f(\mathbf{x}; \theta)$$

There are a few big caveats, in particular:

- The degree of constraint satisfaction depends on the multipliers
- There is no trivial guarantee that constraints are eventually satisfied

...But neither point matters if we just want to inject some expert knowledge in a ML model

A Case Study: Scarce Labels in Predictive Maintenance

Scarce Labels in RUL Predictions

RUL estimation is a major goal for predictive maintenance

However, ground truth for RUL is hard to come by:

- Run-to-failure experiments are time consuming
- They may not be viable for large and complex machines

Typically, **only a few runs** are available

However, data about normal operation may still be abundant

- This may come from test runs, installed machines, etc.
- It looks exactly like the input data for our RUL prediction model
- ...And it will still show sign of component wear

However, **the true RUL value in this case will be unknown**

Can we still take advantage of this data?

Our Baseline Approach

Data Loading and Preparation

We will rely on the NASA C-MAPPS dataset

...Which contains simulated run-to-failure experiments for turbo-fan engines

```
In [4]: data.head()
```

Out[4]:

	src	machine	cycle	p1	p2	p3	s1	s2	s3	s4	...	s13	s14	s15	s16	s17
0	train_FD001	1	1	-0.0007	-0.0004	100.0	518.67	641.82	1589.70	1400.60	...	2388.02	8138.62	8.4195	0.03	392
1	train_FD001	1	2	0.0019	-0.0003	100.0	518.67	642.15	1591.82	1403.14	...	2388.07	8131.49	8.4318	0.03	392
2	train_FD001	1	3	-0.0043	0.0003	100.0	518.67	642.35	1587.99	1404.20	...	2388.03	8133.23	8.4178	0.03	390
3	train_FD001	1	4	0.0007	0.0000	100.0	518.67	642.35	1582.79	1401.87	...	2388.08	8133.83	8.3682	0.03	392
4	train_FD001	1	5	-0.0019	-0.0002	100.0	518.67	642.37	1582.85	1406.22	...	2388.04	8133.80	8.4294	0.03	393

5 rows × 28 columns

- There are four sub-datasets (column `src`)
- Columns `p1-3` represent control parameters
- Columns `s1-21` are sensor readings

Data Loading and Preparation

We will focus on the FD004 dataset (the hardest)

```
In [5]: data_by_src = util.partition_by_field(data, field='src')
dt = data_by_src['train_FD004']
dt[dt_in] = dt[dt_in].astype(np.float32)
```

Then we separate **two sets for training** and one for testing

- The first trainign set will contain finished experiments (supervised)
- ...The second will contain data for still running machines (unsupervised)

```
In [7]: trs_ratio = 0.03 # Supervised experiments / all experiments
tru_ratio = 0.5 # Unsupervised experiments / remaining experiments
trs, tmp = util.split_datasets_by_field(dt, field='machine', fraction=trs_ratio, seed=42)
tru, ts = util.split_datasets_by_field(tmp, field='machine', fraction=tru_ratio, seed=42)

trs_mcn, tru_mcn, ts_mcn = trs['machine'].unique(), tru['machine'].unique(), ts['machine'].unique()
print(f'Num. machine: {len(trs_mcn)} (supervised), {len(tru_mcn)} (unsupervised), {len(ts_mcn)} (test)')
```

```
Num. machine: 7 (supervised), 121 (unsupervised), 121 (test)
```

Data Loading and Preparation

Then we standardize the input data

```
In [9]: sscaler, nscaler = StandardScaler(), MinMaxScaler()
trs_s, tru_s, ts_s = trs.copy(), tru.copy(), ts.copy()
trs_s[dt_in] = sscaler.fit_transform(trs[dt_in])
tru_s[dt_in], ts_s[dt_in] = sscaler.transform(tru[dt_in]), sscaler.transform(ts[dt_in])
trs_s[['rul']] = nscaler.fit_transform(trs[['rul']])
tru_s[['rul']], ts_s[['rul']] = nscaler.transform(tru[['rul']]), nscaler.transform(ts[['rul']])

maxrul = nscaler.data_max_[0]
display(trs_s.head())
```

	src	machine	cycle	p1	p2	p3	s1	s2	s3	s4	...	s13	
1725	train_FD004	467	1	-1.688818	-1.924463	0.445653	1.811018	1.784571	1.676982	1.834240	...	0.445850	0.741
1726	train_FD004	467	2	-0.320795	0.385443	0.445653	0.754417	0.824865	0.604660	0.459056	...	0.445775	-0.158
1727	train_FD004	467	3	-1.688920	-1.925123	0.445653	1.811018	1.768350	1.668955	1.823341	...	0.445477	0.684
1728	train_FD004	467	4	1.184267	0.844852	0.445653	-1.021583	-0.742836	-0.576935	-0.541685	...	0.443309	0.078
1729	train_FD004	467	5	-1.688948	-1.925453	0.445653	1.811018	1.767810	1.726471	1.761244	...	0.445402	0.677

5 rows × 28 columns


Later, we will need the maximum RUL value on the training set

MLP with Scarce Labels

As a baseline, we will train a MLP model **on the supervised data**

We do not split a validation set, given we have scarce data

```
In [13]: nn = util.build_ml_model(input_size=len(dt_in), output_size=1, hidden=[32, 32])  
history = util.train_ml_model(nn, trs_s[dt_in], trs_s['rul'], validation_split=0., epochs=40)  
util.plot_training_history(history, figsize=figsize)
```




Model loss: 0.0235 (training)

Evaluation

Let's inspect the predictions

```
In [14]: trs_pred = nn.predict(trs_s[dt_in]).ravel() * maxrul  
stop = 1400  
util.plot_rul(trs_pred[:stop], trs["rul"].iloc[:stop], same_scale=False, figsize=figsize)
```




- The predictions have a decreasing trend (which is good)
- ...But they are **very noisy** (which is bad)

Evaluation

The behavior on the test data has a similar trend

...And it is similarly noisy

```
In [15]: ts_pred = nn.predict(ts_s[dt_in]).ravel() * maxrul  
util.plot_rul(ts_pred[:stop], ts['rul'].iloc[:stop], same_scale=False, figsize=figsize)
```



Cost Model

The RUL estimator is meant to be used to define a policy

Namely, we stop operations when:

$$f(x, \omega) \leq \theta$$

- Where $f(x, \omega)$ is the estimated output and θ is threshold

Calibrating θ is best done by relying on a cost model

- We assume that operating for a time step generates 1 unit of profit
- ...And that failing loses C units of profits w.r.t. performing maintenance
- We also assume we never stop a machine before a "safe" interval s

Both C and s are calibrated on data in our example:


```
In [16]: failtimes = dt.groupby('machine')['cycle'].max()
         safe_interval, maintenance_cost = failtimes.min(), failtimes.max()
```

Cost Model and Threshold Optimization

We then proceed to choose θ to optimize the cost

```
In [17]: cmodel = util.RULCostModel(maintenance_cost=maintenance_cost, safe_interval=safe_interval)
th_range = np.linspace(-15, 30, 100)
trs_thr = util.optimize_threshold(trs_s['machine'].values, trs_pred, th_range, cmodel, plot=True)
print(f'Optimal threshold for the training set: {trs_thr:.2f}')
```

Optimal threshold for the training set: -2.73



Cost Results

Let's now check the costs on all datasets

```
In [18]: trs_c, trs_f, trs_sl = cmodel.cost(trs_s['machine'].values, trs_pred, trs_thr, return_margin=True)
         ts_c, ts_f, ts_sl = cmodel.cost(ts['machine'].values, ts_pred, trs_thr, return_margin=True)
         print(f'Cost: {trs_c} (supervised), {ts_c} (test)')
```

```
Cost: -406 (supervised), 24566 (test)
```

- The cost for the training set is good (negative)
- ...But that is not the case for the training set

```
In [19]: trs_nm, tru_nm, ts_nm = len(trs_mcn), len(tru_mcn), len(ts_mcn)
         print(f'Avg. fails: {trs_f/trs_nm:.2f} (supervised), {ts_f/ts_nm:.2f} (test)')
         print(f'Avg. slack: {trs_sl/trs_nm:.2f} (supervised), {ts_sl/len(ts_mcn):.2f} (test)')
```

```
Avg. fails: 0.00 (supervised), 0.45 (test)
Avg. slack: 10.57 (supervised), 5.76 (test)
```

- In particular, there is a **very high failure rate on unseen data**

Domain Knowledge as Constraints

Domain Knowledge as Constraints

We know that the RUL **decreases at a fixed rate**

- After 1 time step, the RUL will have decreased by 1 unit
- After 2 time steps, the RUL will have decreased by 2 units and so on

In general, let \hat{x}_i and \hat{x}_j be the i -th and j -th samples **for a given component**

Then we know that:

$$f(\hat{x}_i, \theta) - f(\hat{x}_j, \theta) = j - i \quad \forall i, j = 1..m \text{ s.t. } c_i = c_j$$

- c_i, c_j are the machine for (respectively) sample i and j
- Samples are assumed to be **temporally sorted**
- The left-most terms is the difference between the predicted RULs
- $j - i$ is the difference between the sequential indexes of the two samples
- ...Which by construction should be equal to the RUL difference

Domain Knowledge as Constraints

The relation we identified is a **constraint**

$$f(\hat{x}_i, \theta) - f(\hat{x}_j, \theta) = j - i \quad \forall i, j = 1..m \text{ s.t. } c_i = c_j$$

It represents domain knowledge that should (in principle) hold for our problem

- It is fine to treat this as a **soft constraint**

As a regularization term, we will use:

$$\lambda \left(f(\hat{x}_i, \omega) - f(\hat{x}_j, \omega) - (j - i) \right)^2$$

Using the absolute value (h1 norm) may also work

- The constraint involves pairs of example, i.e. it is a **relational constraint**
- In principle we should consider all pairs, but that may scale poorly

Our Regularizer

We can focus on contiguous pairs, leading to the loss

$$L(\hat{x}, \omega) + \lambda \sum_{\substack{i < j \\ c_i = c_j}} \left(f(\hat{x}_i, \omega) - f(\hat{x}_j, \omega) - (j - i) \right)^2$$

- Where $i < j$ iff j is the next sample for after i for a given machine
- This approach requires a linear (rather than quadratic) number of constraints

It can work with mini-batches

- In this case, $<$ will refer to contiguous samples in the same batch
- ...And of course for the same component

We will now see how to implement this approach

Removing RUL Values

We start by preparing a bit more the unsupervised data

- We remove the end of the unsupervised data sequences
- Then, we replace RUL values with -1 (invalid)
- Finally, we merge supervised and unsupervised data in a single dataset

```
In [20]: tru_s2 = util.rul_cutoff_and_removal(tru_s, cutoff_min=20, cutoff_max=60, seed=42)
tr_s2 = pd.concat((trs_s, tru_s2))
tr_s2.head()
```

Out[20]:

	src	machine	cycle	p1	p2	p3	s1	s2	s3	s4	...	s13	
1725	train_FD004	467	1	-1.688818	-1.924463	0.445653	1.811018	1.784571	1.676982	1.834240	...	0.445850	0.741
1726	train_FD004	467	2	-0.320795	0.385443	0.445653	0.754417	0.824865	0.604660	0.459056	...	0.445775	-0.158
1727	train_FD004	467	3	-1.688920	-1.925123	0.445653	1.811018	1.768350	1.668955	1.823341	...	0.445477	0.684
1728	train_FD004	467	4	1.184267	0.844852	0.445653	-1.021583	-0.742836	-0.576935	-0.541685	...	0.443309	0.078
1729	train_FD004	467	5	-1.688948	-1.925453	0.445653	1.811018	1.767810	1.726471	1.761244	...	0.445402	0.677

5 rows × 28 columns

Generating Batches from the Same Machine

Our regularizer requires to have **sorted samples from the same machine**

The easiest way to ensure we have enough is using a custom `DataGenerator`

```
class SMBatchGenerator(tf.keras.utils.Sequence):  
    def __init__(self, data, in_cols, batch_size, seed=42): ...  
    def __len__(self): ...  
    def __getitem__(self, index): ...  
    def on_epoch_end(self): ...  
    def __build_batches(self): ...
```

- `__len__` is called to know how many batches are left
- `__getitem__` should return one batch
- `on_epoch_end` should take care (e.g.) of shuffling

Generating Batches from the Same Machine

The `__init__` method takes care of the initial setup

```
def __init__(self, data, in_cols, batch_size, seed=42):
    super(SMBatchGenerator).__init__()
    self.data = data
    self.in_cols = in_cols
    self.dpm = split_by_field(data, 'machine')
    self.rng = np.random.default_rng(seed)
    self.batch_size = batch_size
    # Build the first sequence of batches
    self.__build_batches()
```

- We store some fields
- We split the data by machine
- We build a dedicated RNG
- ...And finally we call the custom-made `__build_batches` method

Generating Batches from the Same Machine

The `__build_batches` method prepares the batches for one full epoch

```
def __build_batches(self):
    self.batches, self.machines = [], []
    mcns = list(self.dpm.keys())
    self.rng.shuffle(mcns) # sort the machines at random
    for mcn in mcns: # Loop over all machines
        index = self.dpm[mcn].index # sample indexes for this machine
        ...
        self.rng.shuffle(idx) # shuffle sample indexes for this machine
        bt = idx.reshape(-1, self.batch_size) # split into batches
        bt = np.sort(bt, axis=1) # sort every batch individually
        self.batches.append(bt) # store the batch
        self.machines.append(np.repeat([mcn], len(bt))) # add machine information
    self.batches = np.vstack(self.batches) # concatenate
    self.machines = np.hstack(self.machines)
```

Generating Batches from the Same Machine

We rebuild batches after each epoch

```
def on_epoch_end(self):  
    self.__build_batches()
```

Most of the remaining work is done in the `__getitem__` method:

```
def __getitem__(self, index):  
    idx = self.batches[index]  
    x = self.data[self.in_cols].loc[idx].values  
    y = self.data['rul'].loc[idx].values  
    flags = (y != -1)  
    info = np.vstack((y, flags, idx)).T  
    return x, info
```

- The RUL value is -1 for the unsupervised data: we flag the meaningful RULs
- ...We pack indexes, RUL values, and flags into a single `info` tensor

Custom Training Step

We then enforce the constraints by means of a **custom training step**

```
class CstRULRegressor(keras.Model):  
    def __init__(self, rul_pred, alpha, beta, maxrul): ...  
  
    def train_step(self, data): ...  
  
    def call(self, data): return self.rul_pred(data)  
  
    ...
```

- We use a custom `keras.Model` subclass
- ...And accept an externally built RUL prediction model (`rul_pred`)
- The custom training step is implemented in `train_step`
- The `call` method relies on the external model for RUL prediction

Custom Training Step

In the `__init__` function:

```
def __init__(self, rul_pred, alpha, beta, maxrul):  
    super(CstRULRegressor, self).__init__(input_shape, hidden)  
    # Store the base RUL prediction model  
    self.rul_pred = rul_pred  
    # Weights  
    self.alpha = alpha  
    self.beta = beta  
    self.maxrul = maxrul  
    ...
```

- `beta` is the regularizer weight, `alpha` is a weight for the loss function itself
- We also store the maximum RUL

Custom Training Step

In the custom training step:

```
def train_step(self, data):
    x, info = data
    y_true, flags, idx = info[:, 0:1], info[:, 1:2], info[:, 2:3]
    with tf.GradientTape() as tape:
        y_pred = self(x, training=True) # predictions
        mse = k.mean(flags * k.square(y_pred - y_true)) # MSE loss
        delta_pred = y_pred[1:] - y_pred[:-1] # pred. difference
        delta_rul = -(idx[1:] - idx[:-1]) / self.maxrul # index difference
        deltadiff = delta_pred - delta_rul # difference of differences
        cst = k.mean(k.square(deltadiff)) # regularization term
        loss = self.alpha * mse + self.beta * cst # loss
    ...
```

- We unpack the `info` tensor
- Inside a `GradientTape`, we construct our regularized loss

Custom Training Step

In the custom training step:

```
def train_step(self, data):  
    ...  
    tr_vars = self.trainable_variables  
    grads = tape.gradient(loss, tr_vars) # gradient computation  
  
    self.optimizer.apply_gradients(zip(grads, tr_vars)) # weight update  
  
    ...
```

- We then apply the (Stochastic) Gradient Descent step
- Then we update and return the loss trackers

The SBR Approach

We can now test our SBR approach

```
In [21]: nn_aux = util.build_ml_model(input_size=len(dt_in), output_size=1, hidden=[32, 32])
nn2 = util.CstRULRegressor(rul_pred=nn_aux, alpha=1, beta=5, maxrul=maxrul)
batch_gen = util.CstBatchGenerator(tr_s2, dt_in, batch_size=32)
history = util.train_ml_model(nn2, X=batch_gen, y=None, validation_split=0., epochs=20, verbose=
```

Epoch 1/20

907/907 [=====] - 2s 2ms/step - loss: 0.0416 - mse: 0.0097 - cst: 0.0064

Epoch 2/20

907/907 [=====] - 2s 2ms/step - loss: 0.0129 - mse: 0.0031 - cst: 0.0020

Epoch 3/20

907/907 [=====] - 2s 2ms/step - loss: 0.0129 - mse: 0.0031 - cst: 0.0020

Epoch 4/20

907/907 [=====] - 2s 2ms/step - loss: 0.0129 - mse: 0.0030 - cst: 0.0020

Epoch 5/20

907/907 [=====] - 2s 2ms/step - loss: 0.0126 - mse: 0.0030 - cst: 0.0019

Epoch 6/20

907/907 [=====] - 2s 2ms/step - loss: 0.0124 - mse: 0.0030 - cst: 0.0019


Epoch 7/20

907/907 [=====] - 2s 2ms/step - loss: 0.0125 - mse: 0.0030 - cst: 0.0019

The SBR Approach

...And we can check the training curve

```
In [22]: util.plot_training_history(history, figsize=figsize)
```




Model loss: 0.0118 (training)

Inspecting the Predictions

Let's have a look at the predictions on the supervised data

```
In [23]: trs_pred2 = nn2.predict(trs_s[dt_in]).ravel() * maxrul  
util.plot_rul(trs_pred2[:stop], trs["rul"].iloc[:stop], same_scale=False, figsize=figsize)
```




- The signal is much more stable
- The scale is still off, but we can fix that with a well chosen threshold

Inspecting the Predictions

Then let's do the same for the **unsupervised** data


```
In [24]: tru_pred2 = nn2.predict(tru_s[dt_in]).ravel() * maxrul  
util.plot_rul(tru_pred2[:stop], tru["rul"].iloc[:stop], same_scale=False, figsize=figsize)
```



Inspecting the Predictions

Then let's do the same for the **test** data

```
In [25]: ts_pred2 = nn2.predict(ts_s[dt_in]).ravel() * maxrul  
util.plot_rul(ts_pred2[:stop], ts["rul"].iloc[:stop], same_scale=False, figsize=figsize)
```




The behavior is more stable and consistent than before

Threshold Optimization and Cost Evaluation

We can now optimize the threshold optimization (on the supervised data)

```
In [26]: cmodel = util.RULCostModel(maintenance_cost=maintenance_cost, safe_interval=safe_interval)
th_range2 = np.linspace(50, 150, 200)
trs_thr2 = util.optimize_threshold(trs_s['machine'].values, trs_pred2, th_range2, cmodel, plot=True)
print(f'Optimal threshold for the training set: {trs_thr2:.2f}')
```

Optimal threshold for the training set: 95.23



Threshold Optimization and Cost Evaluation

Finally, we can evaluate the SBR approach in terms of cost

```
In [27]: trs_c2, trs_f2, trs_sl2 = cmodel.cost(trs_s['machine'].values, trs_pred2, trs_thr2, return_margin=True)
tru_c2, tru_f2, tru_sl2 = cmodel.cost(tru_s['machine'].values, tru_pred2, trs_thr2, return_margin=True)
ts_c2, ts_f2, ts_sl2 = cmodel.cost(ts['machine'].values, ts_pred2, trs_thr2, return_margin=True)
print(f'Cost: {trs_c2} (supervised), {tru_c2} (unsupervised), {ts_c2} (test)')
```

```
Cost: -273 (supervised), -7367 (unsupervised), -4665 (test)
```

```
In [28]: print(f'Avg. fails: {trs_f2/len(trs_mcn):.2f} (supervised), {tru_f2/len(tru_mcn):.2f} (unsupervised), {ts_f2/len(ts_mcn):.2f} (test)')
print(f'Avg. slack: {trs_sl2/len(trs_mcn):.2f} (supervised), {tru_sl2/len(tru_mcn):.2f} (unsupervised), {ts_sl2/len(ts_mcn):.2f} (test)')
```

```
Avg. fails: 0.00 (supervised), 0.05 (unsupervised), 0.07 (test)
```

```
Avg. slack: 29.57 (supervised), 34.40 (unsupervised), 32.02 (test)
```

- The number of fails has decreased very significantly
- The slack is still contained

And we did this with **just a handful** of run-to-failure experiments